

# Company Dataset

In [4]:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier
import warnings
warnings.filterwarnings('ignore')
```

In [6]:

```
company = pd.read_csv("C:/Users/HP/Downloads/Company_Data.csv")
company
```

Out[6]:

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	U
0	9.50	138	73	11	276	120	Bad	42	17	
1	11.22	111	48	16	260	83	Good	65	10	
2	10.06	113	35	10	269	80	Medium	59	12	
3	7.40	117	100	4	466	97	Medium	55	14	
4	4.15	141	64	3	340	128	Bad	38	13	
...	...	...	...	...	...	...	...	...	...	...
395	12.57	138	108	17	203	128	Good	33	14	
396	6.14	139	23	3	37	120	Medium	55	11	
397	7.41	162	26	12	368	159	Medium	40	18	
398	5.94	100	79	7	284	95	Bad	50	12	
399	9.71	134	37	0	27	120	Good	49	16	

400 rows × 11 columns



In [7]:

```
company.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sales            400 non-null    float64
1   CompPrice        400 non-null    int64
2   Income           400 non-null    int64
3   Advertising       400 non-null    int64
4   Population        400 non-null    int64
5   Price            400 non-null    int64
6   ShelveLoc        400 non-null    object
7   Age              400 non-null    int64
8   Education         400 non-null    int64
9   Urban            400 non-null    object
10  US                400 non-null    object
dtypes: float64(1), int64(7), object(3)
memory usage: 34.5+ KB
```

In [8]:

```
company.describe()
```

Out[8]:

	Sales	CompPrice	Income	Advertising	Population	Price	Age	E
count	400.000000	400.000000	400.000000	400.000000	400.000000	400.000000	400.000000	400.000000
mean	7.496325	124.975000	68.657500	6.635000	264.840000	115.795000	53.322500	1.000000
std	2.824115	15.334512	27.986037	6.650364	147.376436	23.676664	16.200297	0.000000
min	0.000000	77.000000	21.000000	0.000000	10.000000	24.000000	25.000000	1.000000
25%	5.390000	115.000000	42.750000	0.000000	139.000000	100.000000	39.750000	1.000000
50%	7.490000	125.000000	69.000000	5.000000	272.000000	117.000000	54.500000	1.000000
75%	9.320000	135.000000	91.000000	12.000000	398.500000	131.000000	66.000000	1.000000
max	16.270000	175.000000	120.000000	29.000000	509.000000	191.000000	80.000000	1.000000

In [9]:

```
company.shape
```

Out[9]:

(400, 11)

In [10]:

```
company.isnull().sum()
```

Out[10]:

```
Sales      0
CompPrice  0
Income     0
Advertising 0
Population 0
Price      0
ShelveLoc  0
Age        0
Education  0
Urban      0
US         0
dtype: int64
```

In [11]:

```
from sklearn import preprocessing
label_encoder = preprocessing.LabelEncoder()
```

In [12]:

```
label_encoder
```

Out[12]:

```
LabelEncoder()
```

In [13]:

```
company['ShelveLoc'] = label_encoder.fit_transform(company['ShelveLoc'])
company['Urban'] = label_encoder.fit_transform(company['Urban'])
company['US'] = label_encoder.fit_transform(company['US'])
company.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sales           400 non-null   float64
1   CompPrice       400 non-null   int64
2   Income          400 non-null   int64
3   Advertising     400 non-null   int64
4   Population      400 non-null   int64
5   Price           400 non-null   int64
6   ShelveLoc       400 non-null   int32
7   Age             400 non-null   int64
8   Education       400 non-null   int64
9   Urban           400 non-null   int32
10  US              400 non-null   int32
dtypes: float64(1), int32(3), int64(7)
memory usage: 29.8 KB
```

In [14]:

```
company['value']= company['Sales'].apply(lambda Sales: 'High' if Sales <=7.49 else 'Low')
```

In [15]:

```
company.head()
```

Out[15]:

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	Urban
0	9.50	138	73	11	276	120	0	42	17	
1	11.22	111	48	16	260	83	1	65	10	
2	10.06	113	35	10	269	80	2	59	12	
3	7.40	117	100	4	466	97	2	55	14	
4	4.15	141	64	3	340	128	0	38	13	

In [16]:

```
company.drop(['Sales'],axis=1)
```

Out[16]:

	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	Urban	U
0	138	73	11	276	120	0	42	17	1	
1	111	48	16	260	83	1	65	10	1	
2	113	35	10	269	80	2	59	12	1	
3	117	100	4	466	97	2	55	14	1	
4	141	64	3	340	128	0	38	13	1	
...	...	...	...	...	...	...	...	...	...	...
395	138	108	17	203	128	1	33	14	1	
396	139	23	3	37	120	2	55	11	0	
397	162	26	12	368	159	2	40	18	1	
398	100	79	7	284	95	0	50	12	1	
399	134	37	0	27	120	1	49	16	1	

400 rows × 11 columns

In [17]:

```
company['value'] = label_encoder.fit_transform(company['value'])
```

In [18]:

```
company.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 12 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Sales           400 non-null    float64
 1   CompPrice       400 non-null    int64
 2   Income          400 non-null    int64
 3   Advertising     400 non-null    int64
 4   Population      400 non-null    int64
 5   Price           400 non-null    int64
 6   ShelfLoc        400 non-null    int32
 7   Age             400 non-null    int64
 8   Education       400 non-null    int64
 9   Urban           400 non-null    int32
10   US              400 non-null    int32
11   value           400 non-null    int32
dtypes: float64(1), int32(4), int64(7)
memory usage: 31.4 KB
```

In [19]:

```
x = company.iloc[:,1:11]
y = company.iloc[:,11]
```


In [20]:

```
x
```

Out[20]:

	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	Urban	U
0	138	73	11	276	120	0	42	17	1	
1	111	48	16	260	83	1	65	10	1	
2	113	35	10	269	80	2	59	12	1	
3	117	100	4	466	97	2	55	14	1	
4	141	64	3	340	128	0	38	13	1	
...	...	...	...	...	...	...	...	...	...	...
395	138	108	17	203	128	1	33	14	1	
396	139	23	3	37	120	2	55	11	0	
397	162	26	12	368	159	2	40	18	1	
398	100	79	7	284	95	0	50	12	1	
399	134	37	0	27	120	1	49	16	1	

400 rows × 10 columns



In [21]:

```
y
```

Out[21]:

```
0      1
1      1
2      1
3      0
4      0
..
395    1
396    0
397    0
398    0
399    1
```

Name: value, Length: 400, dtype: int32

In [22]:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.3, random_state=40)
```

In [23]:

```
model=DecisionTreeClassifier(criterion='entropy', max_depth=11)
model.fit(x_train,y_train)
```

Out[23]:

DecisionTreeClassifier(criterion='entropy', max\_depth=11)

In [24]:


```
model.get_n_leaves()
```

Out[24]:

42

In [25]:

```
#plot the Decision Tree
from sklearn import tree
tree.plot_tree(model);
```



In [26]:

```
#Predicting on test data
preds= model.predict(x_test)
#Accuracy on test data
print('Test data Accuracy is:',np.mean(preds==y_test))
```

Test data Accuracy is: 0.6833333333333333

In [27]:

```
#Predicting on train data
predt= model.predict(x_train)
#Accuracy on train data
print('Train data Accuracy is:',np.mean(predt==y_train))
```

Train data Accuracy is: 1.0

In [28]:

```
kfold = KFold(n_splits=10)
cart = DecisionTreeClassifier()
num_trees = 100
model = BaggingClassifier(base_estimator=cart, n_estimators=num_trees)
results = cross_val_score(model, x, y, cv=kfold)
print(results.mean())
```

0.8175000000000001

In [29]:

```
# Random Forest Classification
from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier
num_trees = 100
max_features = 3
kfold = KFold(n_splits=10)
model = RandomForestClassifier(n_estimators=num_trees, max_features=max_features)
results = cross_val_score(model, x, y, cv=kfold)
print(results.mean())
```

0.8175000000000001

In [30]:

```
# AdaBoost Classification
from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import AdaBoostClassifier

num_trees = 10
kfold = KFold(n_splits=10)
model = AdaBoostClassifier(n_estimators=num_trees)
results = cross_val_score(model, x, y, cv=kfold)
print(results.mean())
```

0.7825



In [31]:

```
# Stacking Ensemble for Classification
from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import VotingClassifier
kfold = KFold(n_splits=10)

# create the sub models
estimators = []
model1 = LogisticRegression(max_iter=500)
estimators.append(('logistic', model1))
model2 = DecisionTreeClassifier()
estimators.append(('cart', model2))
model3 = SVC()
estimators.append(('svm', model3))

# create the ensemble model
ensemble = VotingClassifier(estimators)
results = cross_val_score(ensemble, x, y, cv=kfold)
print(results.mean())
```

0.79

## Fraud check dataset

In [32]:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn import tree
from sklearn.metrics import classification_report
import warnings
warnings.filterwarnings('ignore')
```

In [33]:

```
fraud = pd.read_csv("C:/Users/HP/Downloads/Fraud_check.csv")
```

In [34]:

```
fraud.head()
```

Out[34]:

	Undergrad	Marital.Status	Taxable.Income	City.Population	Work.Experience	Urban
0	NO	Single	68833	50047	10	YES
1	YES	Divorced	33700	134075	18	YES
2	NO	Married	36925	160205	30	YES
3	YES	Single	50190	193264	15	YES
4	NO	Married	81002	27533	28	NO

In [35]:

```
fraud.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 600 entries, 0 to 599  
Data columns (total 6 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   Undergrad              600 non-null   object  
1   Marital.Status         600 non-null   object  
2   Taxable.Income         600 non-null   int64  
3   City.Population        600 non-null   int64  
4   Work.Experience        600 non-null   int64  
5   Urban                  600 non-null   object  
dtypes: int64(3), object(3)  
memory usage: 28.2+ KB
```

In [36]:

```
fraud.describe()
```

Out[36]:

	Taxable.Income	City.Population	Work.Experience
count	600.000000	600.000000	600.000000
mean	55208.375000	108747.368333	15.558333
std	26204.827597	49850.075134	8.842147
min	10003.000000	25779.000000	0.000000
25%	32871.500000	66966.750000	8.000000
50%	55074.500000	106493.500000	15.000000
75%	78611.750000	150114.250000	24.000000
max	99619.000000	199778.000000	30.000000

In [37]:

```
fraud.isnull().sum()
```

Out[37]:

```
Undergrad      0
Marital.Status 0
Taxable.Income 0
City.Population 0
Work.Experience 0
Urban          0
dtype: int64
```

In [38]:

```
f1 = fraud.rename({'Marital.Status': 'MarStat', 'Taxable.Income': 'TaxInc', 'City.Population': 'CityPop'})
f1
```

Out[38]:

	Undergrad	MarStat	TaxInc	CityPop	WorkExp	Urban
0	NO	Single	68833	50047	10	YES
1	YES	Divorced	33700	134075	18	YES
2	NO	Married	36925	160205	30	YES
3	YES	Single	50190	193264	15	YES
4	NO	Married	81002	27533	28	NO
...	...	...	...	...	...	...
595	YES	Divorced	76340	39492	7	YES
596	YES	Divorced	69967	55369	2	YES
597	NO	Divorced	47334	154058	0	YES
598	YES	Married	98592	180083	17	NO
599	NO	Divorced	96519	158137	16	NO

600 rows × 6 columns

In [39]:

```
from sklearn import preprocessing
label_encoder = preprocessing.LabelEncoder()
```

In [40]:

```
f1['Undergrad'] = label_encoder.fit_transform(f1['Undergrad'])
f1['MarStat'] = label_encoder.fit_transform(f1['MarStat'])
f1['Urban'] = label_encoder.fit_transform(f1['Urban'])
f1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 600 entries, 0 to 599
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Undergrad    600 non-null    int32
1   MarStat      600 non-null    int32
2   TaxInc       600 non-null    int64
3   CityPop      600 non-null    int64
4   WorkExp      600 non-null    int64
5   Urban        600 non-null    int32
dtypes: int32(3), int64(3)
memory usage: 21.2 KB
```

In [41]:

```
f1['Tax'] = f1['TaxInc'].apply(lambda TaxInc : 'Risky' if TaxInc <= 30000 else 'Good')
```

In [42]:

f1

Out[42]:

	Undergrad	MarStat	TaxInc	CityPop	WorkExp	Urban	Tax
0	0	2	68833	50047	10	1	Good
1	1	0	33700	134075	18	1	Good
2	0	1	36925	160205	30	1	Good
3	1	2	50190	193264	15	1	Good
4	0	1	81002	27533	28	0	Good
...	...	...	...	...	...	...	...
595	1	0	76340	39492	7	1	Good
596	1	0	69967	55369	2	1	Good
597	0	0	47334	154058	0	1	Good
598	1	1	98592	180083	17	0	Good
599	0	0	96519	158137	16	0	Good

600 rows × 7 columns

In [43]:

```
f1=f1.drop(['TaxInc'],axis=1)
f1
```

Out[43]:

	Undergrad	MarStat	CityPop	WorkExp	Urban	Tax
0	0	2	50047	10	1	Good
1	1	0	134075	18	1	Good
2	0	1	160205	30	1	Good
3	1	2	193264	15	1	Good
4	0	1	27533	28	0	Good
...	...	...	...	...	...	...
595	1	0	39492	7	1	Good
596	1	0	55369	2	1	Good
597	0	0	154058	0	1	Good
598	1	1	180083	17	0	Good
599	0	0	158137	16	0	Good

600 rows × 6 columns

In [45]:

```
f1['Tax'] = label_encoder.fit_transform(f1['Tax'])
```

In [46]:

```
f1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 600 entries, 0 to 599
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Undergrad    600 non-null    int32
1   MarStat      600 non-null    int32
2   CityPop      600 non-null    int64
3   WorkExp      600 non-null    int64
4   Urban        600 non-null    int32
5   Tax          600 non-null    int64
dtypes: int32(3), int64(3)
memory usage: 21.2 KB
```

In [47]:

```
x=f1.iloc[:,0:5]
y=f1.iloc[:,5]
```

In [48]:

```
x
```

Out[48]:

	Undergrad	MarStat	CityPop	WorkExp	Urban
0	0	2	50047	10	1
1	1	0	134075	18	1
2	0	1	160205	30	1
3	1	2	193264	15	1
4	0	1	27533	28	0
...	...	...	...	...	...
595	1	0	39492	7	1
596	1	0	55369	2	1
597	0	0	154058	0	1
598	1	1	180083	17	0
599	0	0	158137	16	0

600 rows × 5 columns

In [49]:

```
y
```

Out[49]:

```
0      0
1      0
2      0
3      0
4      0
..
595    0
596    0
597    0
598    0
599    0
Name: Tax, Length: 600, dtype: int64
```

In [50]:

```
y.value_counts()
```

Out[50]:

```
0      476
1      124
Name: Tax, dtype: int64
```

In [51]:

```
#Splitting data into training and testing data set
x_train,x_test,y_train,y_test= train_test_split(x,y, test_size=0.2, random_state=40)
```

In [52]:

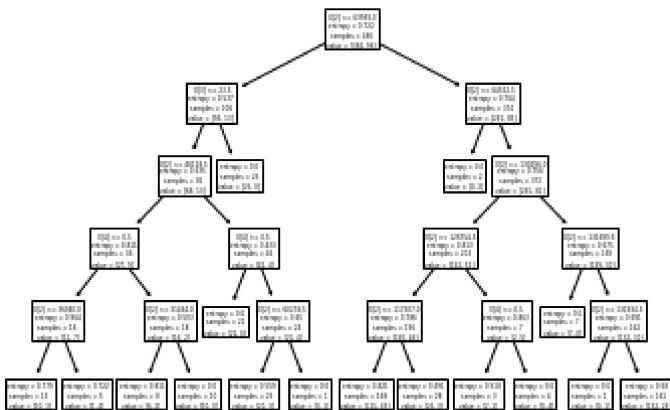
```
model=DecisionTreeClassifier(criterion='entropy', max_depth=5)
model.fit(x_train,y_train)
```

Out[52]:

DecisionTreeClassifier(criterion='entropy', max\_depth=5)

In [53]:

```
#plot the Decision Tree
tree.plot_tree(model);
```



In [54]:

```
model.get_n_leaves()
```

Out[54]:

16

In [55]:

```
#Predicting on test data
preds=model.predict(x_test)
pd.Series(preds).value_counts()
```

Out[55]:

```
0    117
1     3
dtype: int64
```

In [56]:

```
print('Test Data Accuracy:',np.mean(preds==y_test))
```

Test Data Accuracy: 0.775

In [57]:

```
#Predicting on Train Data
predt= model.predict(x_train)
pd.Series(predt).value_counts()
```

Out[57]:

```
0    467
1     13
dtype: int64
```

In [58]:

```
print('Test Data Accuracy:',np.mean(predt==y_train))
```

Test Data Accuracy: 0.8229166666666666

In [59]:

```
#Classification report for y test
print(classification_report(preds,y_test))
```

	precision	recall	f1-score	support
0	0.99	0.78	0.87	117
1	0.07	0.67	0.13	3
accuracy			0.78	120
macro avg	0.53	0.72	0.50	120
weighted avg	0.97	0.78	0.85	120

In [60]:

```
X = f1.iloc[:,0:5]
Y = f1.iloc[:,5]

kfold = KFold(n_splits=10)
cart = DecisionTreeClassifier()
num_trees = 100
model = BaggingClassifier(base_estimator=cart, n_estimators=num_trees)
results = cross_val_score(model, X, Y, cv=kfold)
print(results.mean())
```

0.735

In [61]:

```
# Random Forest Classification
from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier
```



In [62]:

```
X = f1.iloc[:,0:5]
Y = f1.iloc[:,5]
num_trees = 100
max_features = 3
kfold = KFold(n_splits=10)
model = RandomForestClassifier(n_estimators=num_trees, max_features=max_features)
results = cross_val_score(model, X, Y, cv=kfold)
print(results.mean())
```

0.7466666666666667

In [63]:

```
# AdaBoost Classification
from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import AdaBoostClassifier

X = f1.iloc[:,0:5]
Y = f1.iloc[:,5]

num_trees = 10
kfold = KFold(n_splits=10)
model = AdaBoostClassifier(n_estimators=num_trees)
results = cross_val_score(model, X, Y, cv=kfold)
print(results.mean())
```

0.7916666666666666

In [64]:

```
# Stacking Ensemble for Classification
from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import VotingClassifier

X = f1.iloc[:,0:5]
Y = f1.iloc[:,5]
kfold = KFold(n_splits=10)

# create the sub models
estimators = []
model1 = LogisticRegression(max_iter=500)
estimators.append(('logistic', model1))
model2 = DecisionTreeClassifier()
estimators.append(('cart', model2))
model3 = SVC()
estimators.append(('svm', model3))

# create the ensemble model
ensemble = VotingClassifier(estimators)
results = cross_val_score(ensemble, X, Y, cv=kfold)
print(results.mean())
```

0.7933333333333332

In [ ]: